

# Lab2: Parallelizing K-Means with OpenMP

Christos Karagiannis 3544 - Ioannis Sideris 3418 - Athanasios Gerampinis 3466  
University of Thessaly – Department of Electrical and Computer Engineering

## 1 Introduction

This work studies the parallelization of the **K-Means** algorithm using the **OpenMP** framework in C. The goal was to accelerate clustering on large datasets while preserving correctness (identical centers and memberships compared to the sequential version).

The algorithm takes a set of  $N$  points with  $d$  dimensions and computes  $K$  clusters iteratively until convergence (maximum change below a given `threshold`).

## 2 Sequential Implementation

The sequential implementation lives in:

- `seq_kmeans.c` – main compute kernel (`seq_kmeans()`).
- `seq_main.c` – initialization, data I/O, and calling `seq_kmeans()`.

The main function, `seq_kmeans()`, repeats:

1. For each point, call `find_nearest_cluster(...)` (which performs the inner loop over all  $K$  centroids and computes distances).
2. Assign the point to the nearest centroid and update per-cluster sums and counts.
3. Recompute centroids as means of points in each cluster (repeat until the change ratio falls below `threshold` or a max-iteration cap).

These steps are compute-intensive, especially for large  $N$ , and are therefore prime targets for parallelization.

## 3 Parallelization Strategy

Parallelization was applied in the OpenMP version (`omp_kmeans.c`, `omp_main.c`) mainly at the following points:

### 3.1 Distance computation and cluster assignment

The inner double loop:

```
for (i = 0; i < numObjs; i++) {
    int index = find_nearest_cluster(numClusters, numCoords, objects[i], clusters);
    if (membership[i] != index) delta += 1.0;
    membership[i] = index;
```

```

    newClusterSize[index]++;
    for (j = 0; j < numCoords; j++)
        newClusters[index][j] += objects[i][j];
}

```

was parallelized with:

```
#pragma omp parallel for reduction(+:delta) schedule(static)
for (int i = 0; i < numObjs; i++) { ... }
```

so each thread processes a subset of the input points, accelerating the assignment stage.

### 3.2 Updating the new centroids

We handle centroid recomputation in a way that uses *per-thread privatization* with a parallel merge. We keep `schedule(static)` so iterations are split evenly once, with low overhead.

#### 3.2.1 Privatization, merge and normalize

```
#pragma omp parallel shared(ntreads, partialClusterSize, partialClusters) \
                    reduction(+:delta)
{
    #pragma omp single { ntreads = omp_get_num_threads(); }
    #pragma omp for schedule(static)
    for (int i = 0; i < numObjs; i++) {
        int cid = find_nearest_cluster(...);
        localClusterSize[cid]++;
        /* + local sum to partialClusters */
        ...
    }

    #pragma omp parallel for schedule(static)          /* merge */
    for (int c = 0; c < numClusters; c++) {
        int clusterCount = 0; /* + sum partials into newClusters[c][*] */
        ...
    }

    #pragma omp parallel for schedule(static)          /* normalize */
    for (int c = 0; c < numClusters; c++) {
        if (newClusterSize[c] > 0) { float inv = 1.0f / newClusterSize[c]; ... }
        ...
    }
}
```

### 3.3 Scheduling policy

Different `schedule` policies were tested:

- `static` (best performance for uniform load),
- `dynamic` (slightly higher overhead),
- `guided` (no notable improvement).

We finally used `schedule(static)`.

## 4 Optimizations Tried

We experimented with:

- **Parallel sections:** trying to parallelize assignment and recomputation as separate sections did not help, because the latter depends on the former.
- **collapse(2):** applied to the distance loop, but brought no meaningful gains given the small dimensionality (`numCoords = 20`).

## 5 Compiler Flags

Compilation was driven by the `Makefile` with:

```
INCFLAGS      = -I.  
OPTFLAGS     = -ffast-math -DNDEBUG  
OMPFLAGS     = -fopenmp
```

Compiler: Intel(R) oneAPI DPC++/C++ Compiler 2025.2.1 (2025.2.0.20250806).

## 6 Experimental Evaluation

Experiments were run on:

`Image_data/texture17695.bin`

using the following parameters:

- $N = 17695$  objects,
- $d = 20$  coordinates,
- $K = 2000$  clusters.

The sequential runtime was about **3.49 seconds**, while the OpenMP version with 56 threads reached **0.19 seconds**, i.e., **18.36 $\times$**  speedup.

### 6.1 Results

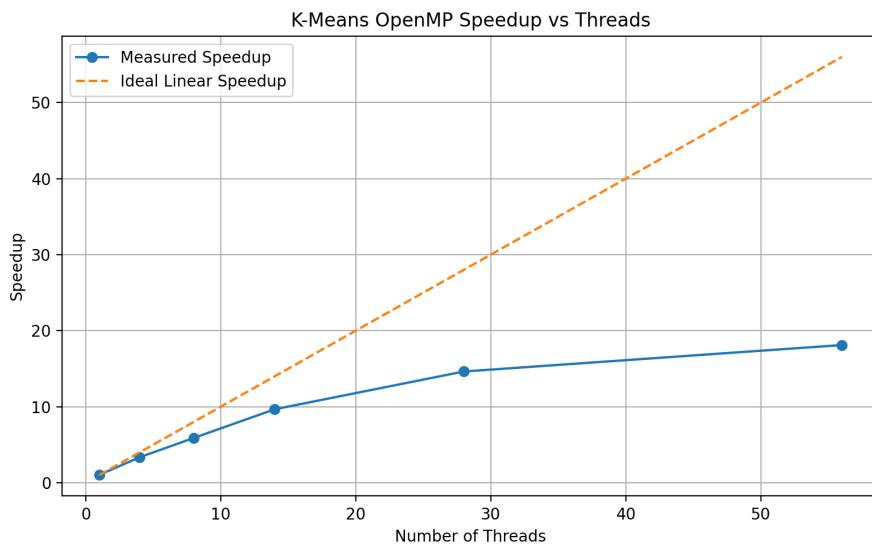


Figure 1: Speedup versus number of threads.

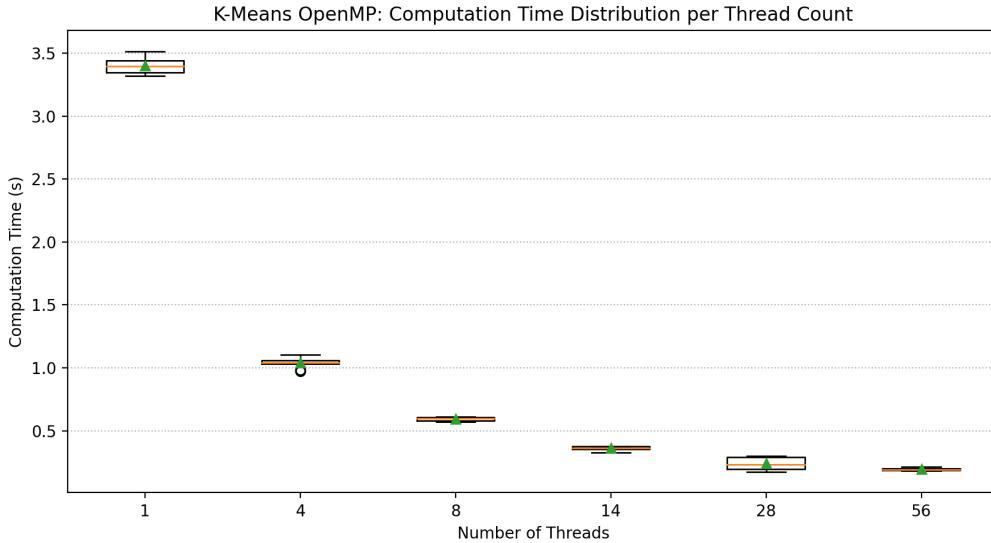


Figure 2: Computation time distribution per thread count.

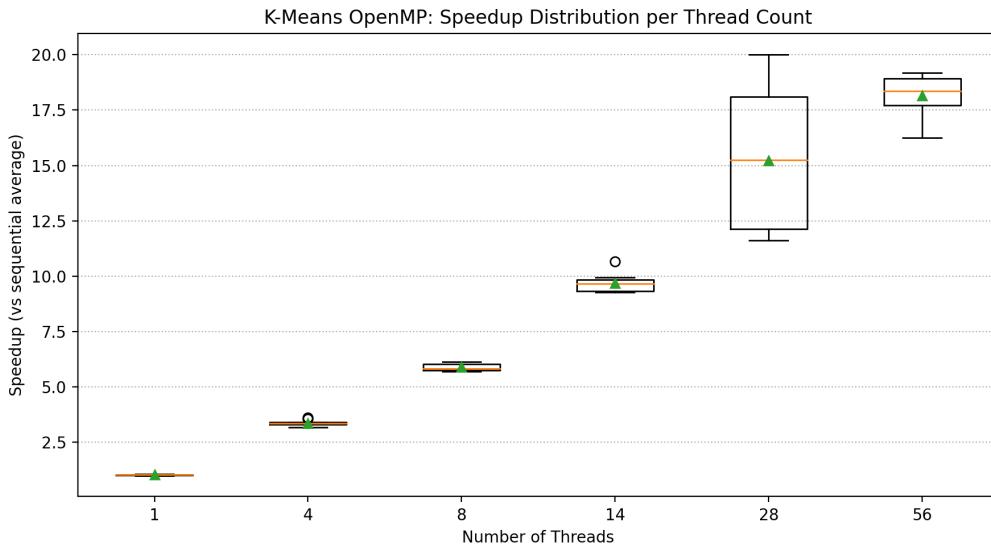


Figure 3: Speedup distribution per thread count.

## 6.2 Discussion

Scaling is good up to roughly 28 threads, close to linear. Beyond that, gains taper off due to:

- memory bandwidth saturation,
- reduced parallelism in the final iterations of K-Means.

The average speedup reaches  $17.6 \times$  at 56 threads, with parallel efficiency around 31%.

## 7 Conclusions

OpenMP parallelization provided significant performance improvements on a multi-core system. Best performance appears when the number of threads matches the physical cores of the CPU.